

**SVKM MITHIBAI COLLEGE OF
SCIENCE,COMMERCE AND ARTS
MUMBAI**

FACIAL EMOTION RECOGNITION

DEEP LEARNING PROJECT

MSC COMPUTER SCIENCE

NAME : VIKRAM VISHNU PANDEY

DIV :F

ROLL NO: 03

YEAR :2024-2025

MENTORED BY : JAYASHREE RAVI MA'AM

Project Introduction

Over the past decade, the Indian education has experienced significant transformation due to the rise of web-based learning services, particularly eLearning platforms. The global eLearning market is projected to grow eightfold over the next five years.

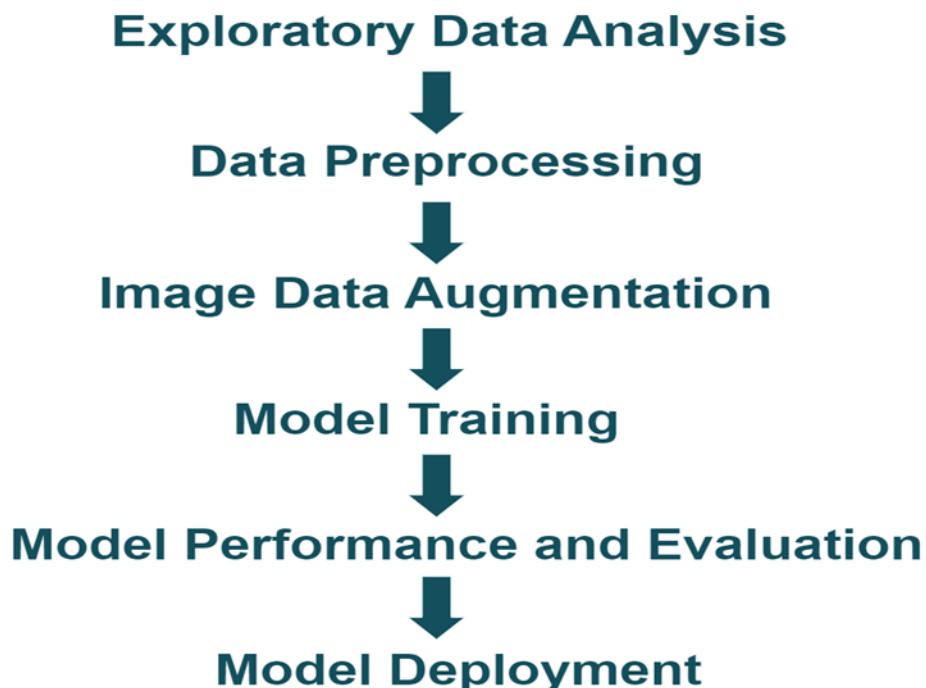
Despite this rapid growth, digital learning faces several challenges compared to traditional brick-and-mortar classrooms. One major challenge is ensuring quality learning for students. While digital platforms may offer superior content quality, assessing whether students are comprehending the material in real-time remains an open question. In physical classrooms, teachers can observe students' facial expressions and emotions, adjusting their teaching pace accordingly and identifying those who need extra attention.

Digital classrooms, often conducted via video conferencing software like Zoom, make it difficult for teachers to monitor the mood and engagement of a medium-sized class (25-50 students). This lack of surveillance can lead to decreased student focus on the content. However, digital platforms have the advantage of leveraging data and machine learning. They generate data in the form of video, audio, and text, which can be analyzed using deep learning algorithms. These **systems not only address the surveillance issue but also eliminate human bias**, translating information from the teacher's observations into quantifiable data that can be **analyzed** and tracked.

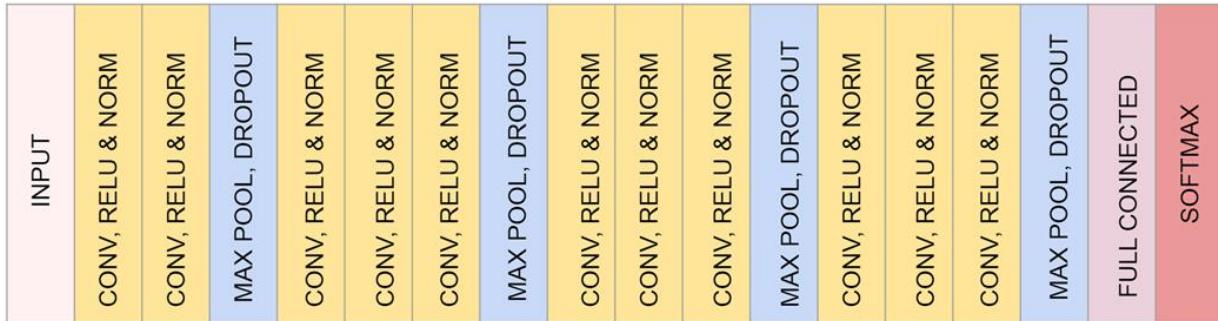
EXPLANATION

1. Here , I will try to solve the above-mentioned challenge by applying deep learning algorithms to live video data.
2. The solution to this problem is by recognizing facial emotions.
3. This is a few shot learning live face emotion detection system. The model should be able to real-time
4. identify the emotions of students in a live class

Flow of the project



Model Architecture:



1. Exploratory Data Analysis (EDA)

- **Objective:** To Understand the dataset's structure, distribution, and potential issues.
- **Steps:**
 - **Data Visualization:** Plot the distribution of different emotions.
 - **Summary Statistics:** Calculate mean, median, and standard deviation for pixel values.
 - **Missing Values:** Identify and handle any missing or corrupted data.
 - **Class Imbalance:** Check for imbalanced classes and plan for handling them.

2. Data Preprocessing

- **Objective:** To Prepare the data for model training.
- **Steps:**

- **Face Detection:** Use algorithms like MTCNN to detect faces in images.
- **Normalization:** Scale pixel values to a range (e.g., 0 to 1).
- **Resizing:** Resize images to a consistent size (e.g., 48x48 pixels).
- **Data Augmentation:** Apply transformations to increase data diversity.

3. Image Data Augmentation

- **Objective:** Enhance the training dataset by creating variations of existing images.
- **Techniques:**
 - **Rotation:** Rotate images by small angles.
 - **Flipping:** Flip images horizontally or vertically.
 - **Scaling:** Zoom in or out on images.
 - **Translation:** Shift images horizontally or vertically.
 - **Brightness Adjustment:** Change the brightness of images.
 -

4. Model Training

- **Objective:** Train a model to recognize facial emotions.
- **Steps:**
 - **Model Selection:** Choose a suitable model architecture
 - **Loss Function:** Use a loss function like categorical cross-entropy.
 - **Optimizer:** Select an optimizer like Adam or SGD.
 - **Training:** Train the model on the preprocessed and augmented data.
 - **Validation:** Validate the model on a separate validation set to monitor performance.

5. Model Performance Evaluation

- **Objective:** Assess the model's performance on unseen data.
- **Metrics:**
 - **Accuracy:** Overall correctness of the model.
 - **Precision, Recall, F1-Score:** Evaluate performance for each class.
 - **Confusion Matrix:** Visualize true vs. predicted classifications.

- **ROC Curve: Assess the trade-off between true positive and false positive rates.**

6. Model Deployment

- **Objective: Integrate the trained model into a live system.**
- **Steps:**
 - **Real-Time Inference: Implement real-time emotion detection from video streams.**

Code :

Notebook Screenshot along with complete Explanation.....within it below in detail

The screenshot shows a Jupyter Notebook interface running in Google Chrome. The left sidebar displays a file tree with a single item: 'fer2013.zip'. The main notebook area has one cell open, titled 'Face_Emotion_Recognition_v0.ipynb'. The code in the cell is as follows:

```
[4]: !pip install pandas numpy matplotlib seaborn pillow tensorflow keras scikit-learn opencv-python
importing libraries
#datatprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from PIL import Image
import tensorflow
from tensorflow.keras.utils import to_categorical
import os
#model_building
from keras.models import Sequential
from keras.layers import Dense,Dropout,Activation,Flatten,BatchNormalization
from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D
from tensorflow.keras.optimizers import Adam, SGD
from keras.callbacks import EarlyStopping
#metric
from sklearn.metrics import classification_report
#real time
import cv2
import warnings
warnings.filterwarnings('ignore')
warnings.simplefilter('ignore')

Requirement already satisfied: pandas in c:\python312\lib\site-packages (2.2.2)
Requirement already satisfied: numpy in c:\python312\lib\site-packages (1.26.4)
Requirement already satisfied: opencv-python in c:\python312\lib\site-packages (4.5.3.1)
Requirement already satisfied: tensorflow in c:\python312\lib\site-packages (2.9.0)
Requirement already satisfied: tensorflow-estimator in c:\python312\lib\site-packages (2.9.0)
Requirement already satisfied: keras in c:\python312\lib\site-packages (2.8.0)
Requirement already satisfied: scikit-learn in c:\python312\lib\site-packages (1.2.2)
Requirement already satisfied: matplotlib in c:\python312\lib\site-packages (3.4.3)
Requirement already satisfied: pillow in c:\python312\lib\site-packages (8.2.0)
Requirement already satisfied: seaborn in c:\python312\lib\site-packages (0.11.2)
Requirement already satisfied: tensorflow-metadata in c:\python312\lib\site-packages (1.9.0)
Requirement already satisfied: tensorflow-quantum in c:\python312\lib\site-packages (0.11.0)
Requirement already satisfied: tensorflow-serve-api in c:\python312\lib\site-packages (0.1.0)
Requirement already satisfied: tensorflow-servemetadata in c:\python312\lib\site-packages (0.1.0)
Requirement already satisfied: tensorflow-quantum-keras in c:\python312\lib\site-packages (0.1.0)
Requirement already satisfied: tensorflow-quantum-optimizer in c:\python312\lib\site-packages (0.1.0)
```

The status bar at the bottom indicates the kernel is 'Python 3 (ipykernel)' and the mode is 'Command'. The date and time shown are '23-09-2024 01:04'.

dataset - JupyterLab

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caet... Time Complexity - ... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_v0.ipynb

Python 3 (ipykernel)

[8]: #check for null values
data.isnull().sum()

[8]: emotion 0
pixels 0
Usage 0
dtype: int64

There are no null values in the dataset.

[9]: #lets check the value counts in usage
data['Usage'].value_counts()

[9]: Usage
Training 28709
PublicTest 3589
PrivateTest 3589
Name: count, dtype: int64

[10]: #lets check the value counts in emotion column
data['emotion'].value_counts()

[10]: emotion
3 8989
6 6198
4 6077
2 5121
0 4953
5 4002
1 547
Name: count, dtype: int64

Mode: Command Ln 87, Col 27 Face_Emotion_Recognition_Vithika_Karan.ipynb 1

NEP - CAN Live

dataset - JupyterLab

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caet... Time Complexity - ... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_v0.ipynb

Python 3 (ipykernel)

[8]: #check for null values
data.isnull().sum()

[8]: PublicTest 3589
PrivateTest 3589
Name: count, dtype: int64

[10]: #lets check the value counts in emotion column
data['emotion'].value_counts()

[10]: emotion
3 8989
6 6198
4 6077
2 5121
0 4953
5 4002
1 547
Name: count, dtype: int64

There are counts for seven emotions. Let's try to explore the images for these classes.

####Exploratory Data Analysis

[11]: import numpy as np
import matplotlib.pyplot as plt

def show_img(df_row, pixel_string_col, class_col):
 ...
 This function takes in pixel data in the form of strings and the respective class of the picture.
 It preprocesses the data and returns an image matrix of 48x48 along with the class label,
 which can be easily plotted to visualize the image.

Mode: Command Ln 87, Col 27 Face_Emotion_Recognition_Vithika_Karan.ipynb 1

NEP - CAN Live

```
####Exploratory Data Analysis

[11]: import numpy as np
       import matplotlib.pyplot as plt

def show_img(df_row, pixel_string_col, class_col):
    ...

    This function takes in pixel data in the form of strings and the respective class of the picture.
    It preprocesses the data and returns an image matrix of 48x48 along with the class label,
    which can be easily plotted to visualize the image.

    Parameters:
    df_row: A row of a dataframe with columns containing pixels as a string and the corresponding class label.
    pixel_string_col: Column name containing the pixels (dtype: string).
    class_col: Column name containing the class label (dtype: string).
    ...

    # Get the pixel data and class label from the dataframe row
    pixels = df_row[pixel_string_col]
    label = df_row[class_col]

    # Convert pixel string to a list of integers
    pic = np.array(pixels.split(), dtype=np.uint8)

    # Reshape the pixel list to a 48x48 array
    pic = pic.reshape(48, 48)

    # Optional: Convert grayscale to RGB by duplicating across channels
```

```
# Convert the pixel list to a 48x48 array
pic = np.array(pic, dtype=np.uint8)

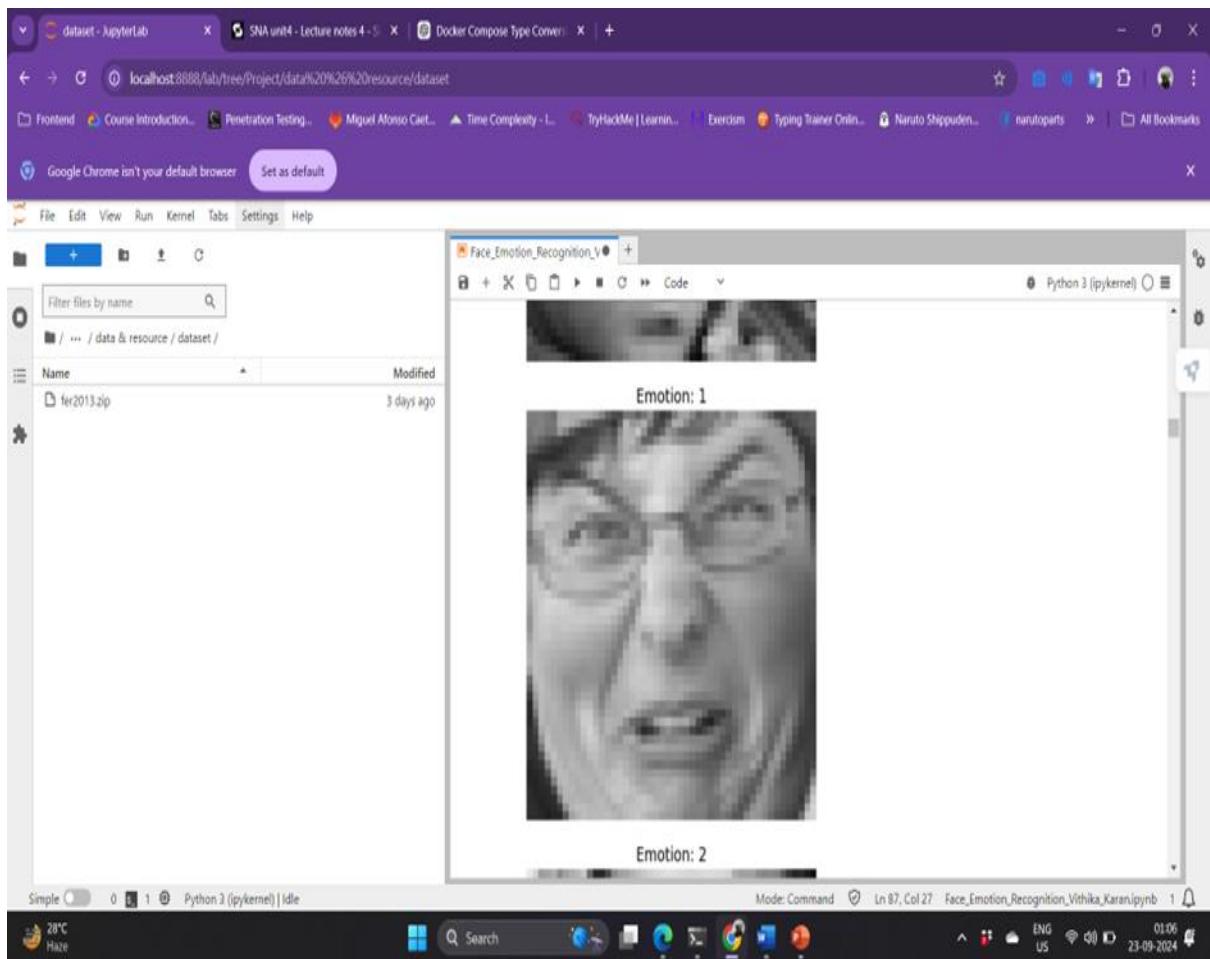
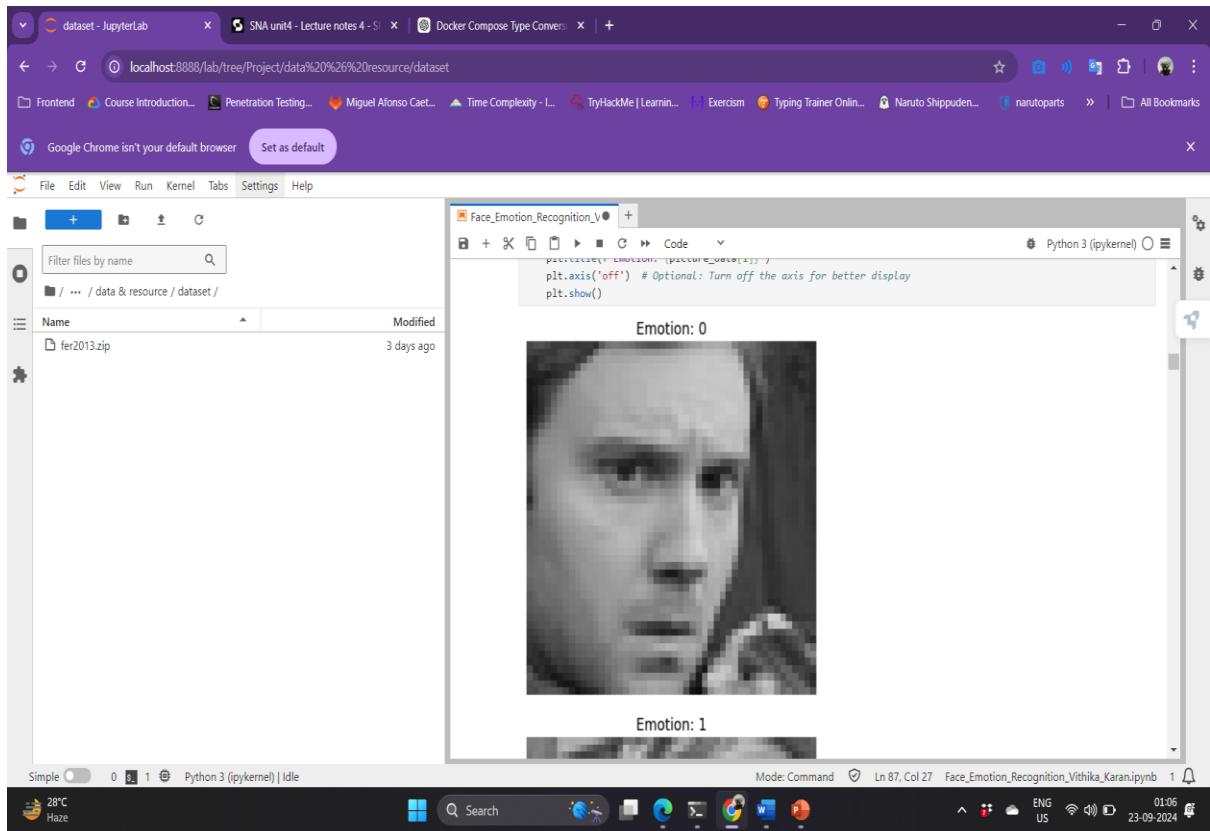
# Reshape the pixel list to a 48x48 array
pic = pic.reshape(48, 48)

# Optional: Convert grayscale to RGB by duplicating across channels
image = np.zeros((48, 48, 3), dtype=np.uint8)
image[:, :, 0] = pic # Red channel
image[:, :, 1] = pic # Green channel
image[:, :, 2] = pic # Blue channel

# Return the image (3-channel) and the label as a tuple
return image, label

# Iterate over each emotion and plot the corresponding pixels
for emotion in range(0, 7): # Assuming emotion labels range from 0 to 6
    picture = data[data['emotion'] == emotion].iloc[0] # Select the first picture for each emotion
    picture_data = show_img(picture, 'pixels', 'emotion')

    # Display the image
    plt.imshow(picture_data[0]) # Show the RGB image
    plt.title(f"Emotion: {picture_data[1]}")
    plt.axis('off') # Optional: Turn off the axis for better display
    plt.show()
```



dataset - JupyterLab SNA unit4 – Lecture notes 4 - S... Docker Compose Type Convers...

localhost:8888/lab/tree/Project/data%20%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caet... Time Complexity - I... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_V... Python 3 (ipykernel)

Data Description

The data consists of grayscale images of faces at a resolution of 48x48 pixels. The faces have been automatically registered such that they are more or less centred in each image and take up around the same amount of area.

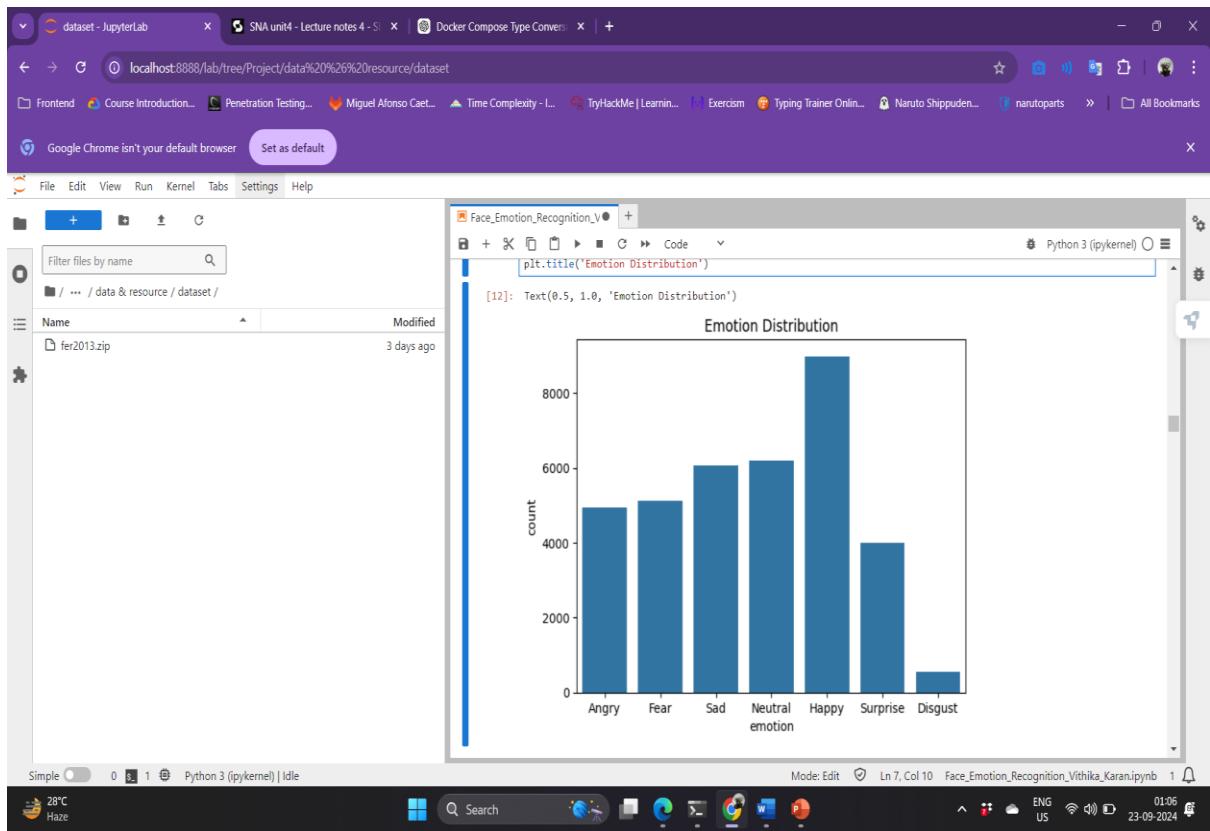
The goal is to categorize each face into one of seven categories based on the emotion expressed in the facial expression (0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral). There are 28,709 examples in the training set and 3,589 examples in the public test set.

```
[12]: #plotting a bar graph of the emotions
# create dict of emotions
emo_dict = {0:'Angry', 1:'Disgust', 2:'Fear', 3:'Happy', 4:'Sad', 5:'Surprise',6:'Neutral'}
#creating a copy of data with actual categorial emotions
df1 = data.copy()
df1['emotion'] = df1['emotion'].map(emo_dict)
#plotting
sns.countplot(x=df1['emotion'])
plt.title("Emotion Distribution")
```

Text(0.5, 1.0, 'Emotion Distribution')

Emotion Distribution

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Edit Ln 7, Col 10 Face_Emotion_Recognition_Virthika_Karan.ipynb 1 01:06 28°C Haze Search ENG US 23-09-2024



The number of happy pictures are the highest and disgusted pictures lowest in the dataset.

```
####Data Pre-processing
```

```
[13]: #splitting the data into train, validation and test set
train_data = data[data['Usage']=='Training']
val_data = data[data['Usage']=='PublicTest']
test_data = data[data['Usage']=='PrivateTest']
print("The shape of training set is: {}, \nThe shape of validation set is: {}, \nThe shape of test set is: {}".format(train_data.shape, val_data.shape, test_data.shape))

The shape of training set is: (28709, 3),
The shape of validation set is: (3589, 3),
The shape of test set is: (3589, 3)

[14]: #resetting index and dropping usage
for elem in [train_data,val_data,test_data]:
    elem.reset_index(inplace=True)
    elem.drop('Usage',axis=1,inplace=True)

[15]: # Data preparation for test
def process_pixels(pixels):
```

The method flow from directory for Image Data Augmentation expects that images belonging to different classes are present in different folders but are inside the same parent folder. Let's create that.

```
[16]: #main_data_directory_path
```

```
[16]: #main data directory path
path = "Project/data & resource"

[17]: import os

def create_dir(path, class_list):
    ...
    The function takes in the path and list of the classes to
    create directories for different classes.

    Parameters:
    path: The path where train and validation directories need to be created.
    class_list: The list of labels in the dataset.
    ...

    # Create train and validation directories
    train_path = os.path.join(path, 'train')
    val_path = os.path.join(path, 'valid')

    try:
        os.mkdir(train_path)
        os.mkdir(val_path)
    except FileExistsError:
        print(f"Directory {train_path} or {val_path} already exists.")

    for data_path, cat in (train_path: 'train-', val_path: 'valid-').items():
        for label in class_list:
            label_dir = os.path.join(data_path, cat + str(label))
            try:
                os.mkdir(label_dir)
            except FileExistsError:
                print(f"Directory {label_dir} already exists.")

[18]: #creating directories
create_dir(path,[0,1,2,3,4,5,6])
```

```
[18]: #creating directories
create_dir(path,[0,1,2,3,4,5,6])

[19]: Directory Project/data & resource\train or Project\data & resource\valid already exists.
Directory Project\data & resource\train\train-0 already exists.
Directory Project\data & resource\train\train-1 already exists.
Directory Project\data & resource\train\train-2 already exists.
Directory Project\data & resource\train\train-3 already exists.
Directory Project\data & resource\train\train-4 already exists.
Directory Project\data & resource\train\train-5 already exists.
Directory Project\data & resource\train\train-6 already exists.
Directory Project\data & resource\valid\valid-0 already exists.
Directory Project\data & resource\valid\valid-1 already exists.
Directory Project\data & resource\valid\valid-2 already exists.
Directory Project\data & resource\valid\valid-3 already exists.
Directory Project\data & resource\valid\valid-4 already exists.
Directory Project\data & resource\valid\valid-5 already exists.
Directory Project\data & resource\valid\valid-6 already exists.

[19]: train_path = "Project/data & resource/train"
val_path = "Project/data & resource/valid"

[20]: #saving Images
def save_imgs(df,df_path,pixel_col,class_col,class_list,prefix):
    ...
    This function takes in the dataframes and
```

The image shows two side-by-side Jupyter Notebook interfaces running on a Windows operating system.

Top Notebook (Python 3 (ipykernel)):

```
[19]: train_path = "Project/data & resource/train"
val_path = "Project/data & resource/valid"

[20]: #saving images
def save_imgs(df,df_path,pixel_col,class_col,class_list,prefix):
    '''This function takes in the dataframes and
    creates images and saves images in directories.
    Parameters:
    df: Dataframe that needs to be converted.
    df_path: Path to the directory (datatype:string)
    Example- If the training dataframe is fed, df_path should be the path
    to the train directory created.
    pixel_col: Name of the column containing pixels in string object
    class_col: Name of the column for data labels
    prefix: train- for training set, valid- for validation set '''

    for i in range(len(df)):
        pixel_string = df[pixel_col][i]
        pixels = list(map(int, pixel_string.split()))

        matrix = np.array(pixels).reshape(48,48).astype(np.uint8)
        img = Image.fromarray(matrix)
        for label in class_list:
            if str(df[class_col][i]) in prefix + str(label):
                img.save(df_path +'/' +prefix + str(label) + '/' + prefix + str(label) + '-' + str(i) + '.png')
            else:
                continue
```

Bottom Notebook (Python 3 (ipykernel)):

```
[21]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
def img_data_gen(train_path, val_path, target_size, batch_size, color_mode, class_mode):
    ...
    This function generates augmented training data and rescales validation images, given the paths.

    Parameters:
    train_path: Path to the training directory of images.
    val_path: Path to the validation directory of images.
    target_size: Target size for images (example: (48, 48)).
    batch_size: Number of images per batch.
    color_mode: Color mode, e.g., 'grayscale' or 'rgb'.
    class_mode: Classification mode, e.g., 'binary' or 'categorical'.

    Returns:
    train_generator: Generator for the training set.
    val_generator: Generator for the validation set.
    ...

    # Image data generator for training with augmentation
    train_datagen = ImageDataGenerator(
        rescale=1./255,
        rotation_range=10,
        horizontal_flip=True,
        width_shift_range=0.1,
        height_shift_range=0.1,
```

dataset - JupyterLab SNA unit4 – Lecture notes 4 - Slides Docker Compose Type Conversion

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caetano Time Complexity - I... TryHackMe | Learning... Exercism Typing Trainer Online... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_v1.ipynb

```
# Image data generator for training with augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=10,
    horizontal_flip=True,
    width_shift_range=0.1,
    height_shift_range=0.1,
    fill_mode='nearest'
)

# Image data generator for validation (no augmentation)
val_datagen = ImageDataGenerator(rescale=1./255)

# Training data generator
train_generator = train_datagen.flow_from_directory(
    train_path,
    target_size=target_size,
    batch_size=batch_size,
    color_mode=color_mode,
    seed=42,
    shuffle=True,
    class_mode=class_mode
)

# Validation data generator
val_generator = val_datagen.flow_from_directory(
    val_path,
    target_size=target_size,
    batch_size=batch_size,
    color_mode=color_mode
)

# Validation data generator
val_generator = val_datagen.flow_from_directory(
    val_path,
    target_size=target_size,
    batch_size=batch_size,
    color_mode=color_mode,
    seed=42,
    shuffle=True, # You can set shuffle=False if you want consistent validation
    class_mode=class_mode
)

return train_generator, val_generator

[23]: #train images
save_imgs(train_data,train_path,'pixels','emotion',[0,1,2,3,4,5,6],'train-')
#validation images
save_imgs(val_data,val_path,'pixels','emotion',[0,1,2,3,4,5,6],'valid-')

[24]: #image generators
train_gen, val_gen = img_data_gen(train_path, val_path, target_size=(48,48), batch_size=64, color_mode='grayscale', class_mode='categorical')
Found 28709 images belonging to 7 classes.
Found 3589 images belonging to 7 classes.

####Model Building
```

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Edit Ln 2, Col 36 Face_Emotion_Recognition_Vithika_Karan.ipynb 1 01:08 23-09-2024 USD/INR -0.16%

dataset - JupyterLab SNA unit4 – Lecture notes 4 - Slides Docker Compose Type Conversion

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caetano Time Complexity - I... TryHackMe | Learning... Exercism Typing Trainer Online... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_v1.ipynb

```
# Validation data generator
val_generator = val_datagen.flow_from_directory(
    val_path,
    target_size=target_size,
    batch_size=batch_size,
    color_mode=color_mode,
    seed=42,
    shuffle=True, # You can set shuffle=False if you want consistent validation
    class_mode=class_mode
)

return train_generator, val_generator

[23]: #train images
save_imgs(train_data,train_path,'pixels','emotion',[0,1,2,3,4,5,6],'train-')
#validation images
save_imgs(val_data,val_path,'pixels','emotion',[0,1,2,3,4,5,6],'valid-')

[24]: #image generators
train_gen, val_gen = img_data_gen(train_path, val_path, target_size=(48,48), batch_size=64, color_mode='grayscale', class_mode='categorical')
Found 28709 images belonging to 7 classes.
Found 3589 images belonging to 7 classes.

####Model Building
```

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Edit Ln 2, Col 36 Face_Emotion_Recognition_Vithika_Karan.ipynb 1 01:08 23-09-2024 USD/INR -0.16%

The screenshot shows a Jupyter Notebook interface running in a browser window. The notebook is titled "Face_Emotion_Recognition_Vithika_Karan.ipynb".

The code in the notebook is as follows:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Convolution2D, BatchNormalization, MaxPooling2D, Dropout, Flatten, Dense
from tensorflow.keras.optimizers import Adam

model = Sequential()
num_classes = 7

# 1st Block
model.add(Convolution2D(64, kernel_size=3, activation='relu', padding='same', input_shape=(48, 48, 1)))
model.add(BatchNormalization())
model.add(Convolution2D(64, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

# 2nd Block
model.add(Convolution2D(128, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Convolution2D(128, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Convolution2D(128, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())

# 3rd Block
model.add(Convolution2D(256, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Convolution2D(256, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Convolution2D(256, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.3))

# 4th Block
model.add(Convolution2D(256, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Convolution2D(256, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Convolution2D(256, kernel_size=3, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.3))

# Flatten and Fully Connected layers
model.add(Flatten())

# Optional Dense Layers
model.add(Dense(512, activation='relu'))
```

The screenshot shows a Jupyter Notebook interface with a Python kernel. The code cell contains the following:

```

model.add(Dropout(0.3))

# Flatten and Fully Connected Layers
model.add(Flatten())

# Optional Dense Layers
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))

model.add(Dense(num_classes, activation='softmax'))

# Summary of the model
model.summary()

# Compile the model
model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

The output cell displays the model summary:

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 48, 48, 64)	640
batch_normalization_11 (BatchNormalization)	(None, 48, 48, 64)	256
conv2d_12 (Conv2D)	(None, 48, 48, 64)	36,928
batch_normalization_12 (BatchNormalization)	(None, 48, 48, 64)	256

The screenshot shows a Jupyter Notebook interface with a Python kernel. The code cell contains the same code as the first one:

```

model.add(Dropout(0.3))

# Flatten and Fully Connected Layers
model.add(Flatten())

# Optional Dense Layers
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))

model.add(Dense(num_classes, activation='softmax'))

# Summary of the model
model.summary()

# Compile the model
model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

The output cell displays the model summary:

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 48, 48, 64)	640
batch_normalization_11 (BatchNormalization)	(None, 48, 48, 64)	256
conv2d_12 (Conv2D)	(None, 48, 48, 64)	36,928
batch_normalization_12 (BatchNormalization)	(None, 48, 48, 64)	256
max_pooling2d_4 (MaxPooling2D)	(None, 24, 24, 64)	0
dropout_5 (Dropout)	(None, 24, 24, 64)	0
conv2d_13 (Conv2D)	(None, 24, 24, 128)	73,856
batch_normalization_13 (BatchNormalization)	(None, 24, 24, 128)	512
conv2d_14 (Conv2D)	(None, 24, 24, 128)	147,584
batch_normalization_14 (BatchNormalization)	(None, 24, 24, 128)	512
conv2d_15 (Conv2D)	(None, 24, 24, 128)	147,584
batch_normalization_15 (BatchNormalization)	(None, 24, 24, 128)	512

dataset - JupyterLab

SNA unit4 - Lecture notes 4 - S | Docker Compose Type Convers | +

localhost:8888/lab/tree/Project/data%20%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alonso Caet... Time Complexity - L... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_vX

		Python 3 (ipykernel)	
batch_normalization_14	(None, 24, 24, 128)	512	
conv2d_15	(Conv2D)	(None, 24, 24, 128)	147,584
batch_normalization_15	(BatchNormalization)	(None, 24, 24, 128)	512
max_pooling2d_5	(MaxPooling2D)	(None, 12, 12, 128)	0
dropout_6	(Dropout)	(None, 12, 12, 128)	0
conv2d_16	(Conv2D)	(None, 12, 12, 256)	295,168
batch_normalization_16	(BatchNormalization)	(None, 12, 12, 256)	1,024
conv2d_17	(Conv2D)	(None, 12, 12, 256)	590,080
batch_normalization_17	(BatchNormalization)	(None, 12, 12, 256)	1,024
conv2d_18	(Conv2D)	(None, 12, 12, 256)	590,080
batch_normalization_18	(BatchNormalization)	(None, 12, 12, 256)	1,024
max_pooling2d_6	(MaxPooling2D)	(None, 6, 6, 256)	0
dropout_7	(Dropout)	(None, 6, 6, 256)	0
conv2d_19	(Conv2D)	(None, 6, 6, 256)	590,080

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Edit Ln 2, Col 36 Face_Emotion_Recognition_Vithika_Karan.ipynb 1 28°C Haze Search ENG US 01:09 23-09-2024

dataset - JupyterLab

SNA unit4 - Lecture notes 4 - S | Docker Compose Type Convers | +

localhost:8888/lab/tree/Project/data%20%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alonso Caet... Time Complexity - L... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_vX

		Python 3 (ipykernel)	
conv2d_19	(Conv2D)	(None, 6, 6, 256)	590,080
batch_normalization_19	(BatchNormalization)	(None, 6, 6, 256)	1,024
conv2d_20	(Conv2D)	(None, 6, 6, 256)	590,080
batch_normalization_20	(BatchNormalization)	(None, 6, 6, 256)	1,024
conv2d_21	(Conv2D)	(None, 6, 6, 256)	590,080
batch_normalization_21	(BatchNormalization)	(None, 6, 6, 256)	1,024
max_pooling2d_7	(MaxPooling2D)	(None, 3, 3, 256)	0
dropout_8	(Dropout)	(None, 3, 3, 256)	0
flatten_1	(Flatten)	(None, 2304)	0
dense_2	(Dense)	(None, 512)	1,180,160
dropout_9	(Dropout)	(None, 512)	0
dense_3	(Dense)	(None, 7)	3,591

Total params: 4,844,103 (18.48 MB)
Trainable params: 4,840,087 (18.46 MB)
Non-trainable params: 4,096 (16.00 KB)

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Edit Ln 2, Col 36 Face_Emotion_Recognition_Vithika_Karan.ipynb 1 CAD/INR -0.24% Search ENG US 01:10 23-09-2024

dataset - JupyterLab

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alonso Caet... Time Complexity - ... TryHackMe | Learnin... Exercism Typing Trainer Online... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_Vithika_Karan.ipynb

```
[27]: #model compiling
optimizer = Adam(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy', optimizer=optimizer, metrics=['accuracy'])

[28]: #training
earlystop = EarlyStopping(monitor='val_loss', min_delta=0, patience = 10, verbose=1, restore_best_weights=True)

epochs = 40
#fitting
history=model.fit(train_gen,
                    steps_per_epoch=len(train_data)//64,
                    epochs=epochs,
                    batch_size=32,
                    callbacks=[earlystop],
                    verbose=1,
                    validation_data=val_gen,
                    validation_steps=len(val_data)//64)
```

Epoch 1/40
448/448 1528s 3s/step - accuracy: 0.1964 - loss: 2.5658 - val_accuracy: 0.2215 - val_loss: 1.8468
Epoch 2/40
448/448 3s 495us/step - accuracy: 0.2500 - loss: 1.8236 - val_accuracy: 0.2000 - val_loss: 1.5871
Epoch 3/40
448/448 1288s 3s/step - accuracy: 0.2344 - loss: 1.8205 - val_accuracy: 0.2634 - val_loss: 1.7924
Epoch 4/40

CAD/INR -0.24% 0 1 Python 3 (ipykernel) | Idle Q Search ENG US 01:10 23-09-2024

dataset - JupyterLab

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alonso Caet... Time Complexity - ... TryHackMe | Learnin... Exercism Typing Trainer Online... Naruto Shippuden... narutoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_Vithika_Karan.ipynb

```
Epoch 4/40  
448/448 3s 221us/step - accuracy: 0.1719 - loss: 1.9128 - val_accuracy: 0.2000 - val_loss: 2.2499  
Epoch 5/40  
448/448 1265s 3s/step - accuracy: 0.2559 - loss: 1.7897 - val_accuracy: 0.3050 - val_loss: 1.6980  
Epoch 6/40  
448/448 3s 229us/step - accuracy: 0.2188 - loss: 1.7289 - val_accuracy: 0.2000 - val_loss: 1.6660  
Epoch 7/40  
448/448 1270s 3s/step - accuracy: 0.2933 - loss: 1.7429 - val_accuracy: 0.3064 - val_loss: 1.8483  
Epoch 8/40  
448/448 3s 230us/step - accuracy: 0.3594 - loss: 1.6105 - val_accuracy: 0.2000 - val_loss: 2.2703  
Epoch 9/40  
448/448 1550s 3s/step - accuracy: 0.3251 - loss: 1.6790 - val_accuracy: 0.3898 - val_loss: 1.5846  
Epoch 10/40  
448/448 4s 347us/step - accuracy: 0.4062 - loss: 1.6248 - val_accuracy: 0.4000 - val_loss: 1.9182  
Epoch 11/40  
448/448 1868s 4s/step - accuracy: 0.3610 - loss: 1.6241 - val_accuracy: 0.3870 - val_loss: 1.6470  
Epoch 12/40  
448/448 4s 341us/step - accuracy: 0.4219 - loss: 1.5431 - val_accuracy: 0.4000 - val_loss: 1.9136  
Epoch 13/40  
448/448 1843s 4s/step - accuracy: 0.3872 - loss: 1.5674 - val_accuracy: 0.4277 - val_loss: 1.4630  
Epoch 14/40
```

CAD/INR -0.24% 0 1 Python 3 (ipykernel) | Idle Q Search ENG US 01:10 23-09-2024

The screenshot shows a Jupyter Notebook interface running in Google Chrome. The top navigation bar has tabs for 'dataset - JupyterLab', 'SNA unit4 - Lecture notes 4 - Slides', and 'Docker Compose Type Convers'. The address bar shows 'localhost:8888/lab/tree/Project/data%20%26%20resource/dataset'. Below the address bar is a bookmarks bar with links like 'Frontend', 'Course Introduction...', 'Penetration Testing...', 'Miguel Afonso Caetano', 'Time Complexity - I...', 'TryHackMe | Learnin...', 'Exercism', 'Typing Trainer Online...', 'Naruto Shippuden...', 'narutoparts', and 'All Bookmarks'. A message 'Google Chrome isn't your default browser' is displayed with a 'Set as default' button.

The main area contains a file browser on the left showing a directory structure with 'fer2013.zip' and a code cell on the right displaying the training logs for a Face Emotion Recognition model. The code cell output is as follows:

```
Face_Emotion_Recognition_V1.ipynb
Python 3 (ipykernel) | Idle
1.4030
Epoch 14/40
448/448 - 5s 408us/step - accuracy: 0.5156 - loss: 1.4368 - val_accuracy: 0.4000 - val_loss: 1.3920
Epoch 15/40
448/448 - 1863s 4s/step - accuracy: 0.4088 - loss: 1.5196 - val_accuracy: 0.4361 - val_loss: 1.4583
Epoch 16/40
448/448 - 4s 403us/step - accuracy: 0.2969 - loss: 1.8360 - val_accuracy: 0.4000 - val_loss: 1.8812
Epoch 17/40
448/448 - 1799s 4s/step - accuracy: 0.4437 - loss: 1.4460 - val_accuracy: 0.4548 - val_loss: 1.4822
Epoch 18/40
448/448 - 4s 416us/step - accuracy: 0.3750 - loss: 1.6130 - val_accuracy: 0.6000 - val_loss: 1.1916
Epoch 19/40
448/448 - 1775s 4s/step - accuracy: 0.4717 - loss: 1.3901 - val_accuracy: 0.5128 - val_loss: 1.2708
Epoch 20/40
448/448 - 4s 417us/step - accuracy: 0.3906 - loss: 1.4183 - val_accuracy: 0.8000 - val_loss: 0.8731
Epoch 21/40
448/448 - 1586s 3s/step - accuracy: 0.4787 - loss: 1.3555 - val_accuracy: 0.5262 - val_loss: 1.2248
Epoch 22/40
448/448 - 3s 230us/step - accuracy: 0.5938 - loss: 1.1069 - val_accuracy: 0.6000 - val_loss: 1.6133
Epoch 23/40
448/448 - 1289s 3s/step - accuracy: 0.4971 - loss: 1.3083 - val_accuracy: 0.5273 - val_loss: 1.2334
```

This screenshot is identical to the one above, showing the same Jupyter Notebook interface, file browser, and code cell output for the Face Emotion Recognition model. The logs show the training progress from epoch 14 to 23, with various metrics like accuracy, loss, and validation metrics displayed.

dataset - JupyterLab SNA unit4 - Lecture notes 4 - S... Docker Compose Type Conver... +

localhost:8888/lab/tree/Project%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caet... Time Complexity - I... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... nanotoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_vX

```
Code Python 3 (ipykernel) 3 days ago
448/448 1299s 3s/step - accuracy: 0.5695 - loss: 1.1441 - val_accuracy: 0.5806 - val_loss: 0.6906
Epoch 36/40
448/448 3s 260us/step - accuracy: 0.5781 - loss: 1.0125 - val_accuracy: 0.8000 - val_loss: 1.1211
Epoch 37/40
448/448 1651s 4s/step - accuracy: 0.5740 - loss: 1.1319 - val_accuracy: 0.5781 - val_loss: 1.1991
Epoch 38/40
448/448 4s 380us/step - accuracy: 0.6406 - loss: 1.0494 - val_accuracy: 0.2000 - val_loss: 1.7910
Epoch 39/40
448/448 1736s 4s/step - accuracy: 0.5842 - loss: 1.1157 - val_accuracy: 0.5806 - val_loss: 1.1065
Epoch 40/40
448/448 4s 333us/step - accuracy: 0.5625 - loss: 1.0326 - val_accuracy: 0.2000 - val_loss: 1.7910
Epoch 40: early stopping
Restoring model weights from the end of the best epoch: 30.

Restored the best model weights, which was for the epoch 45 with the following result:
• loss: 0.7075 - accuracy: 0.7381 - val_loss: 0.9825 - val_accuracy: 0.6643
```

####Model Evaluation

```
[29]: #saving the history of the model in dataframe
model_df = pd.DataFrame(model.history.history)
```

Python 3 (ipykernel) | Idle

CAD/INR -0.24%

dataset - JupyterLab SNA unit4 - Lecture notes 4 - S... Docker Compose Type Conver... +

localhost:8888/lab/tree/Project%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alfonso Caet... Time Complexity - I... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... nanotoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_vX

```
Code Python 3 (ipykernel) 3 days ago
####Model Evaluation
```

```
[29]: #saving the history of the model in dataframe
model_df = pd.DataFrame(model.history.history)
```

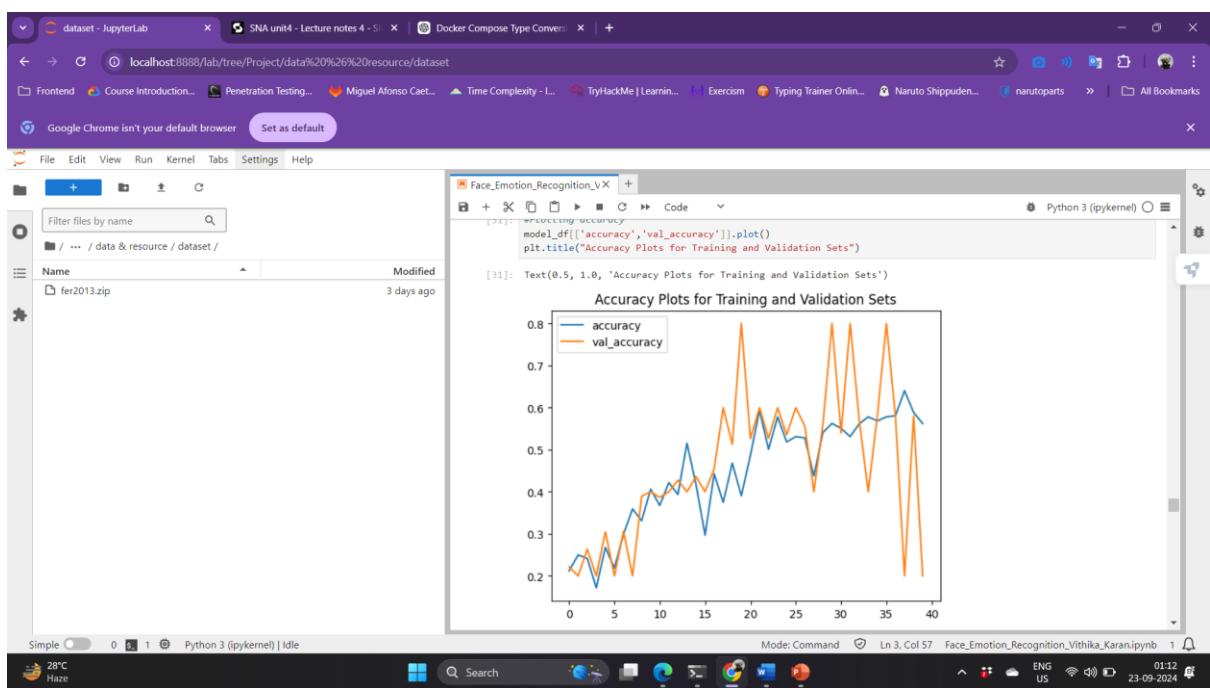
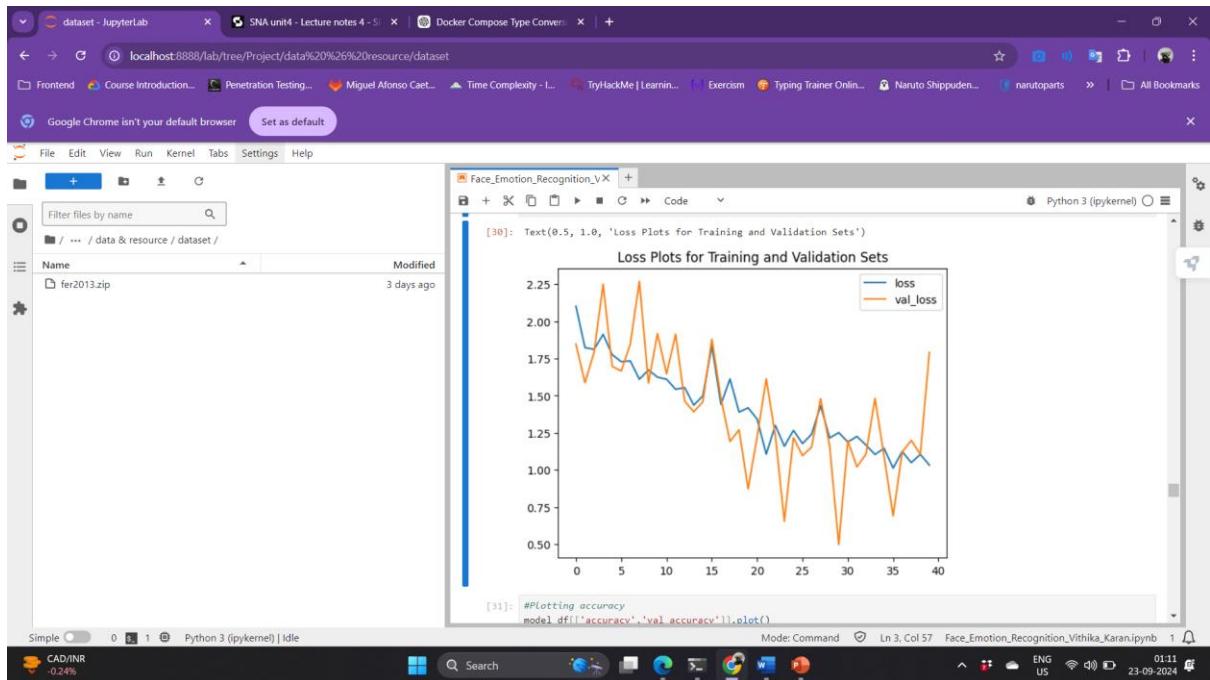
```
[30]: #Plotting loss
model_df[['loss','val_loss']].plot()
plt.title('Loss Plots for Training and Validation Sets')
Text(0.5, 1.0, 'Loss Plots for Training and Validation Sets')
```

Loss Plots for Training and Validation Sets

loss
val_loss

Mode: Command Ln 3, Col 57 Face_Emotion_Recognition_Vithika_Karan.ipynb

0111 CAD/INR -0.24% 23-09-2024



dataset - JupyterLab

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alonso Caet... Time Complexity - I... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... nanotoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_vX

Predictions on Test Data

```
[32]: #test predictions
true_y = np.argmax(y_test, axis=1)
pred_y = np.argmax(model.predict(X_test), axis=1)
print(classification_report(true_y, pred_y))

113/113 66s 584ms/step
precision    recall   f1-score   support
          0       0.54      0.37      0.44     491
          1       0.45      0.09      0.15      55
          2       0.42      0.17      0.24     528
          3       0.73      0.87      0.79     879
          4       0.41      0.49      0.45     594
          5       0.68      0.72      0.70     416
          6       0.48      0.65      0.55     626

accuracy                           0.68
macro avg       0.53      0.48      0.48     3589
weighted avg    0.55      0.57      0.54     3589
```

The predictions on the test set are pretty great with accuracy of 0.68 which implies generalization of the model is good as compared to the training results.

####Save the Model

Simple Python 3 (ipykernel) | Idle Mode: Command Ln 3, Col 57 Face_Emotion_Recognition_Vithika_Karan.ipynb 01:12 23-09-2024

dataset - JupyterLab

localhost:8888/lab/tree/Project/data%20%26%20resource/dataset

Frontend Course Introduction... Penetration Testing... Miguel Alonso Caet... Time Complexity - I... TryHackMe | Learnin... Exercism Typing Trainer Onlin... Naruto Shippuden... nanotoparts All Bookmarks

Google Chrome isn't your default browser Set as default

File Edit View Run Kernel Tabs Settings Help

Face_Emotion_Recognition_vX

The predictions on the test set are pretty great with accuracy of 0.68 which implies generalization of the model is good as compared to the training results.

####Save the Model

```
[39]: #saving model
model.save(path='Emotion Detection Model/emotion_detection.keras')

####Real Time Emotion Detection
```

The following code for real time emotion detection was run on local device and the results are attached along with it.

```
[40]: #Load model
model = tensorflow.keras.models.load_model(path='Emotion Detection Model/emotion_detection.keras')

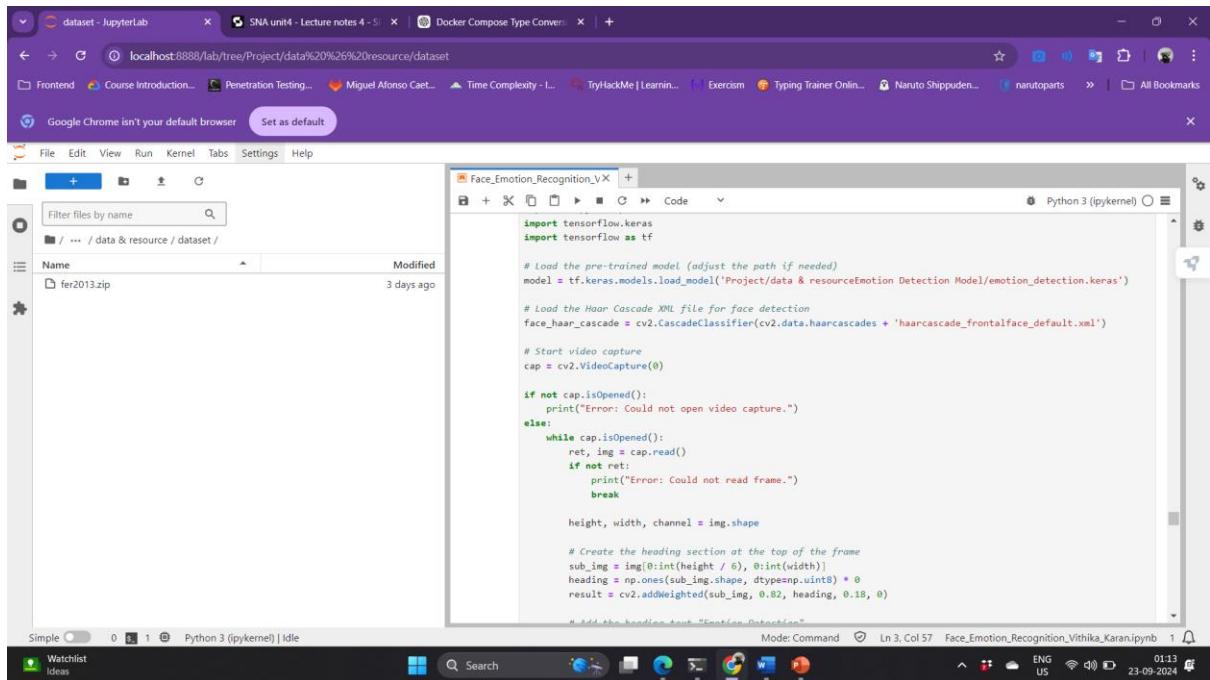
[41]: face_haar_cascade = cv2.CascadeClassifier(path='haarcascade_frontalface_default.xml')

[ ]: import cv2
import numpy as np
import tensorflow.keras
import tensorflow as tf

# Load the pre-trained model (adjust the path if needed)
model = tf.keras.models.load_model('Project/data & resourceEmotion Detection Model/emotion_detection.keras')

# Load the Haar Cascade XML file for face detection.
```

Simple Python 3 (ipykernel) | Idle Mode: Command Ln 3, Col 57 Face_Emotion_Recognition_Vithika_Karan.ipynb 01:13 23-09-2024



```
import tensorflow.keras
import tensorflow as tf

# Load the pre-trained model (adjust the path if needed)
model = tf.keras.models.load_model('Project/data & resource/Motion Detection Model/emotion_recognition.keras')

# Load the Haar Cascade XML file for face detection
face_haar_cascade = cv2.CascadeClassifier(cv2.data.haarcascades + 'haarcascade_frontalface_default.xml')

# Start video capture
cap = cv2.VideoCapture(0)

if not cap.isOpened():
    print("Error: Could not open video capture.")
else:
    while cap.isOpened():
        ret, img = cap.read()
        if not ret:
            print("Error: Could not read frame.")
            break

        height, width, channel = img.shape

        # Create the heading section at the top of the frame
        sub_img = img[0:int(height / 6), 0:int(width)]
        heading = np.ones(sub_img.shape, dtype=np.uint8) * 0
        result = cv2.addWeighted(sub_img, 0.82, heading, 0.18, 0)

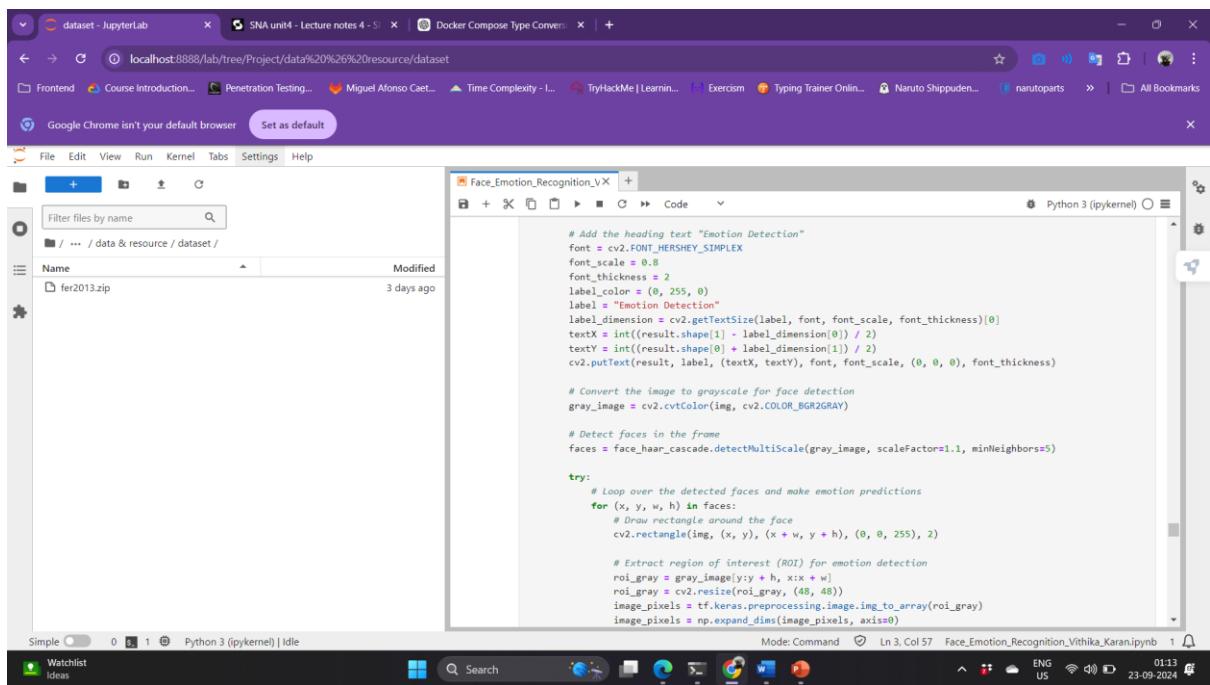
        # Add the heading text "Emotion Detection"
        font = cv2.FONT_HERSHEY_SIMPLEX
        font_scale = 0.8
        font_thickness = 2
        label_color = (0, 255, 0)
        label = "Emotion Detection"
        label_dimension = cv2.getTextSize(label, font, font_scale, font_thickness)[0]
        textx = int((result.shape[1] - label_dimension[0]) / 2)
        texty = int((result.shape[0] + label_dimension[1]) / 2)
        cv2.putText(result, label, (textx, texty), font, font_scale, (0, 0, 0), font_thickness)

        # Convert the image to grayscale for face detection
        gray_image = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # Detect faces in the frame
        faces = face_haar_cascade.detectMultiScale(gray_image, scaleFactor=1.1, minNeighbors=5)

        try:
            # Loop over the detected faces and make emotion predictions
            for (x, y, w, h) in faces:
                # Draw rectangle around the face
                cv2.rectangle(img, (x, y), (x + w, y + h), (0, 0, 255), 2)

                # Extract region of interest (ROI) for emotion detection
                roi_gray = gray_image[y:y + h, x:x + w]
                roi_gray = cv2.resize(roi_gray, (48, 48))
                image_pixels = tf.keras.preprocessing.image.img_to_array(roi_gray)
                image_pixels = np.expand_dims(image_pixels, axis=0)
```



```
                image_pixels = np.expand_dims(image_pixels, axis=0)

                # Predict emotion
                prediction = model.predict(image_pixels)
                emotion_label = np.argmax(prediction)

                # Add text label to the ROI
                cv2.putText(img, emotion_label, (x, y), cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 0, 255), 2)

        except:
            pass

# Release the video capture and close all windows
cap.release()
cv2.destroyAllWindows()
```

The screenshot shows a Jupyter Notebook interface. On the left, there is a file browser window titled 'dataset - JupyterLab' showing a single file 'fer2013.zip'. On the right, a code editor window titled 'Face_Emotion_Recognition_vX' contains the following Python script:

```
cv2.rectangle(img, (x, y), (x + w, y + h), (0, 0, 255), 2)

# Extract region of interest (ROI) for emotion detection
roi_gray = gray[y:y + h, x:x + w]
roi_gray = cv2.resize(roi_gray, (48, 48))
image_pixels = tf.keras.preprocessing.image.img_to_array(roi_gray)
image_pixels = np.expand_dims(image_pixels, axis=0)
image_pixels /= 255 # Normalize pixel values

# Predict the emotion
predictions = model.predict(image_pixels)
max_index = np.argmax(predictions[0])

# Map the prediction to an emotion label
emotion_labels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Sad', 'Surprised', 'Neutral']
emotion_prediction = emotion_labels[max_index]

# Display the predicted emotion on the frame
cv2.putText(img, emotion_prediction, (x, y - 10), font, 0.9, label_color, 2)

except Exception as e:
    print("Error during prediction: (e)")

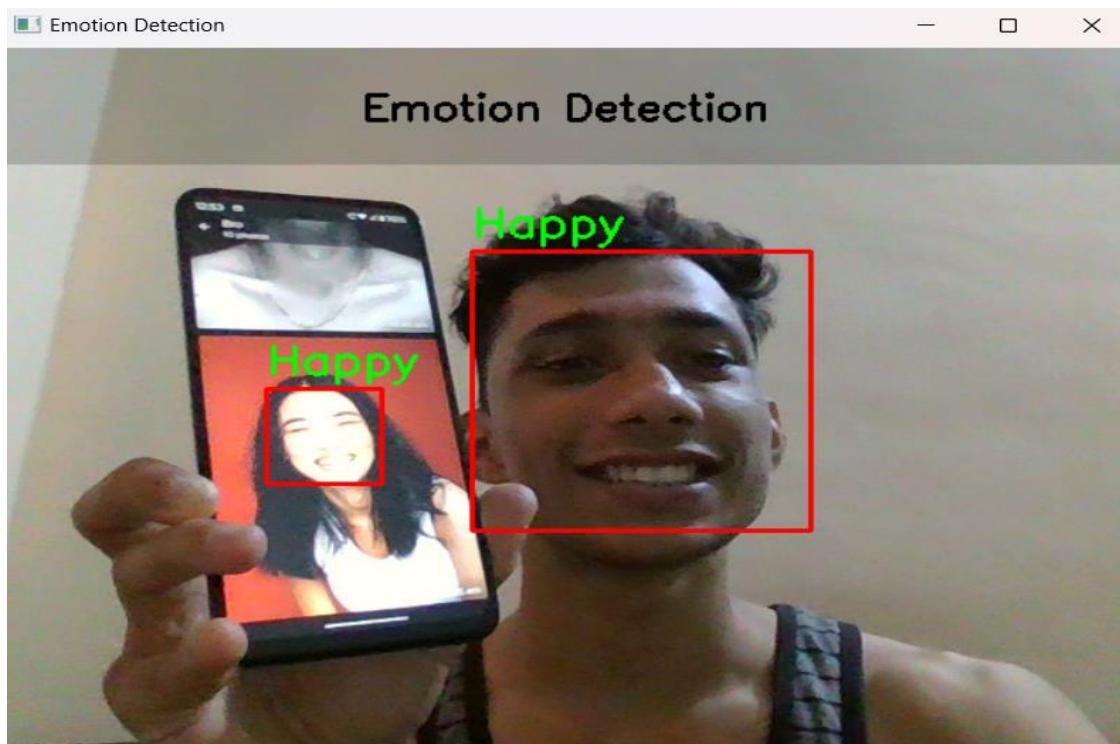
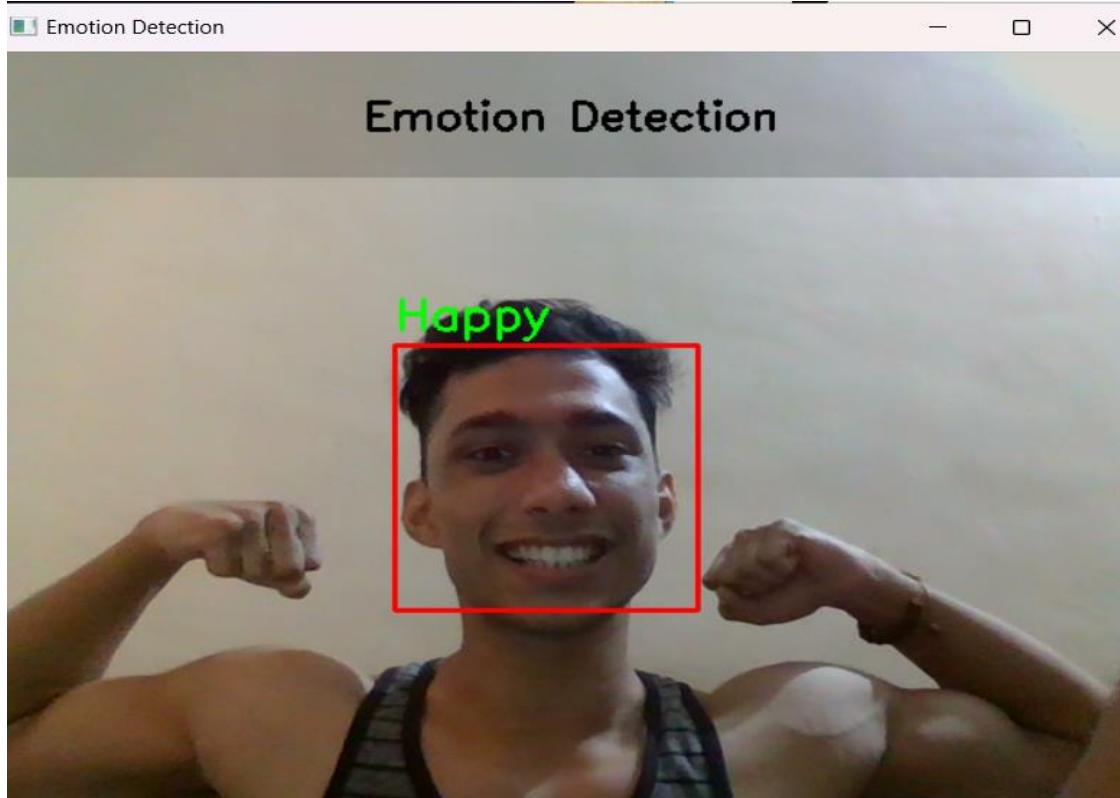
# Overlay the heading on the frame
img[0:int(height / 6), 0:int(width)] = result

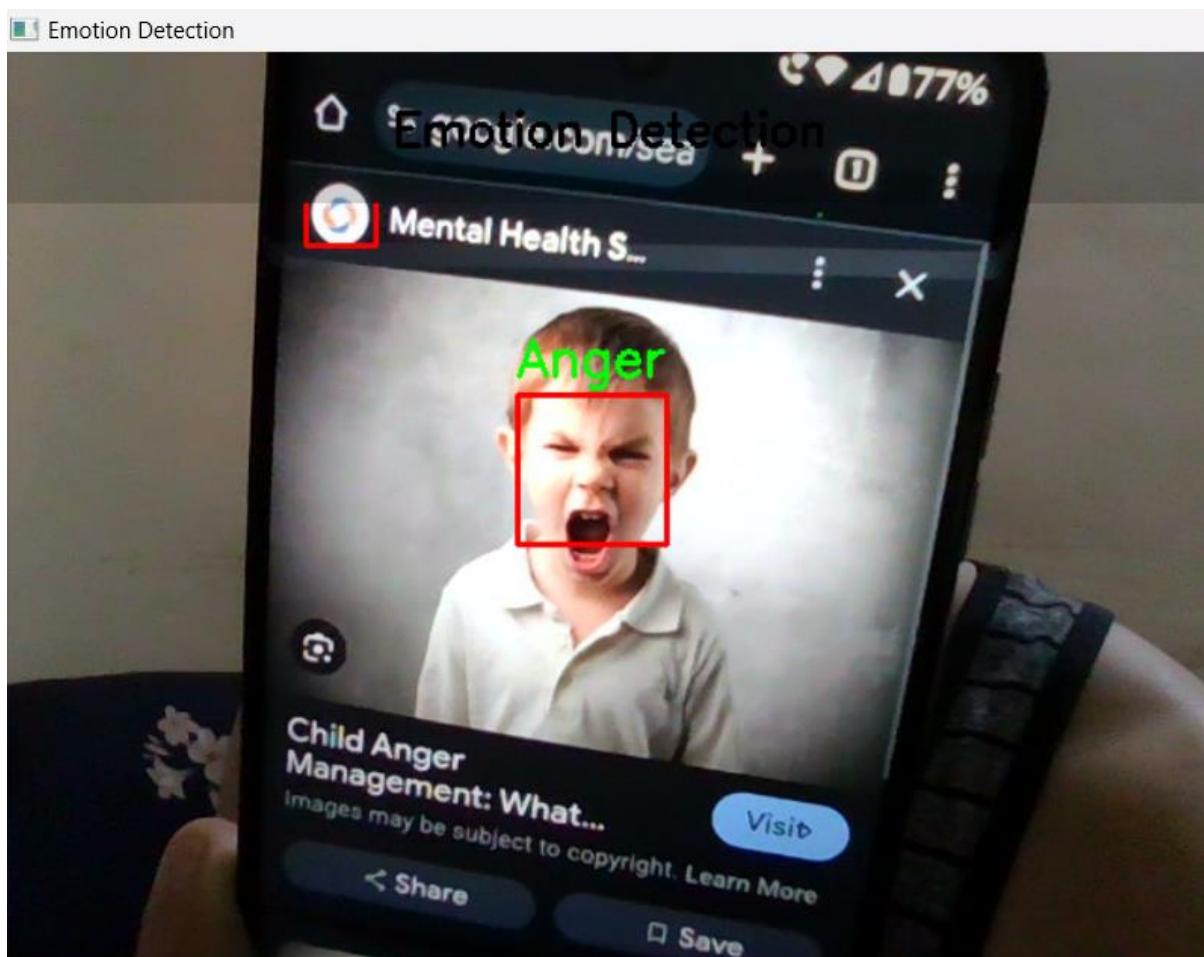
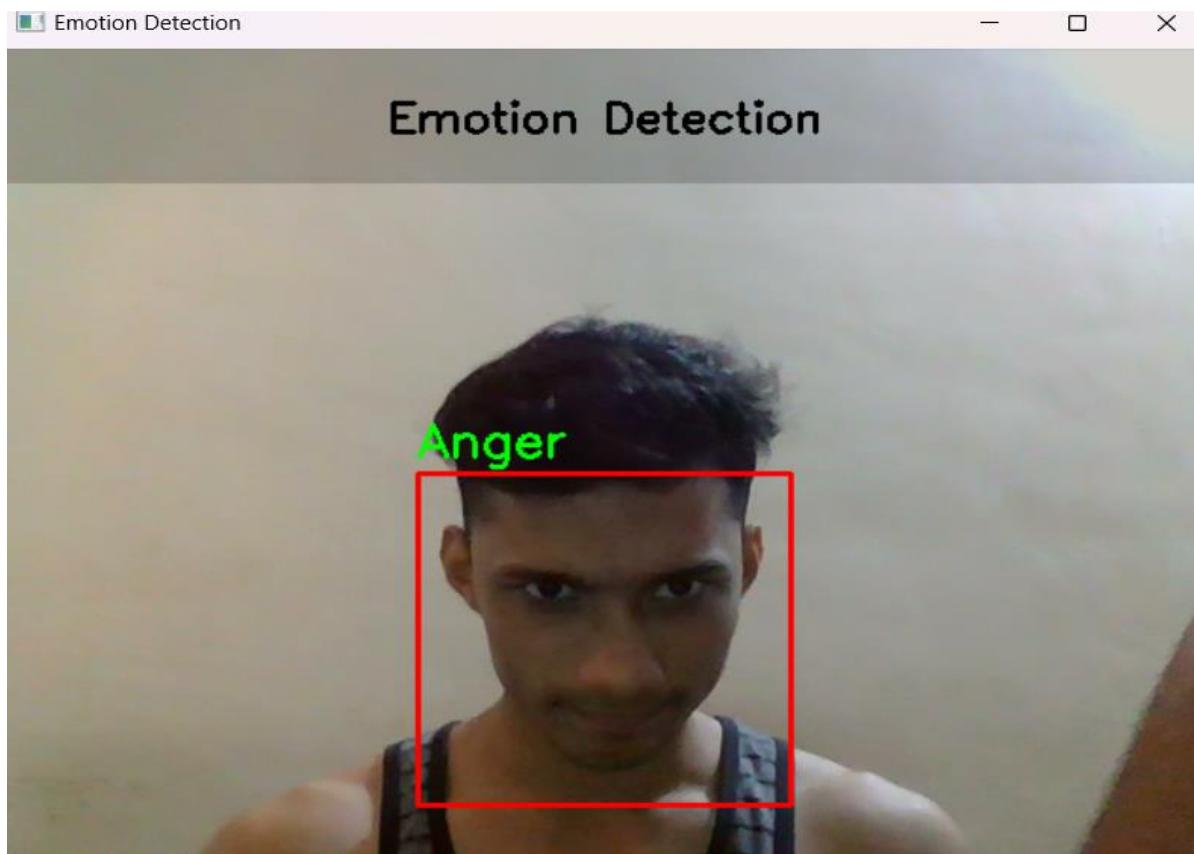
# Show the frame with detected emotion
cv2.imshow('Emotion Detection', img)
```

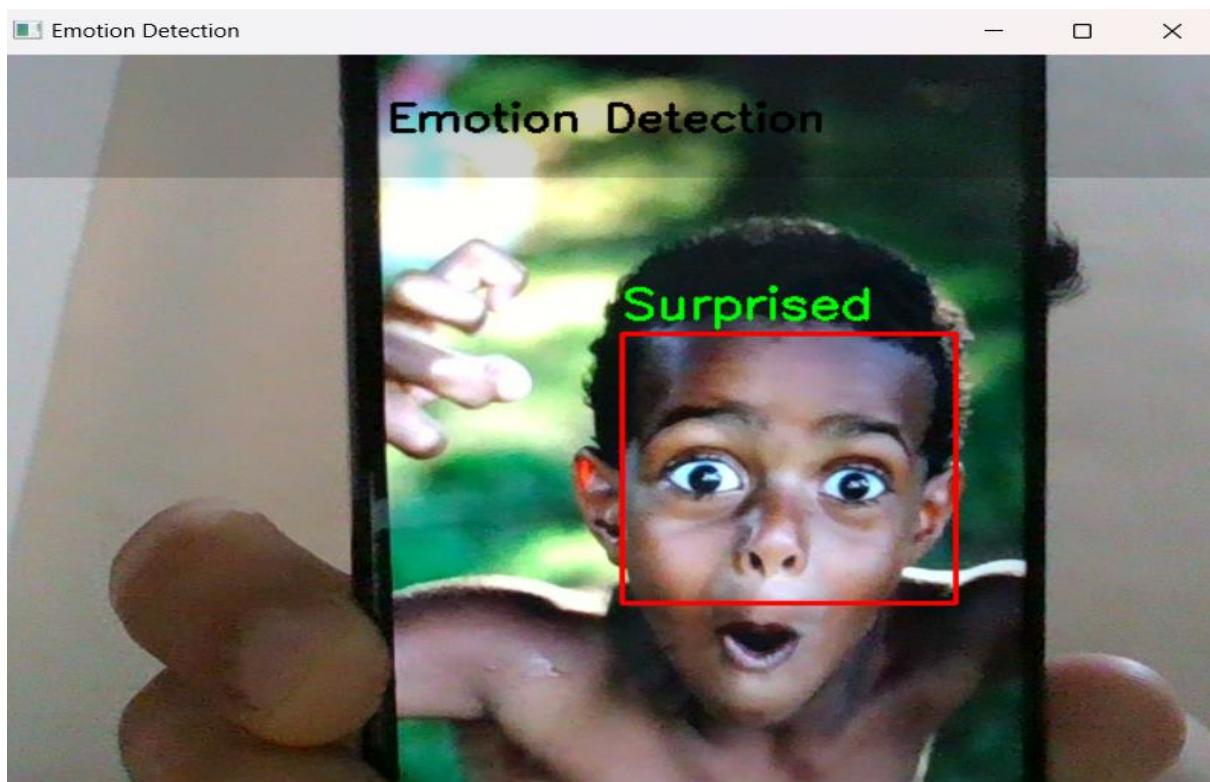
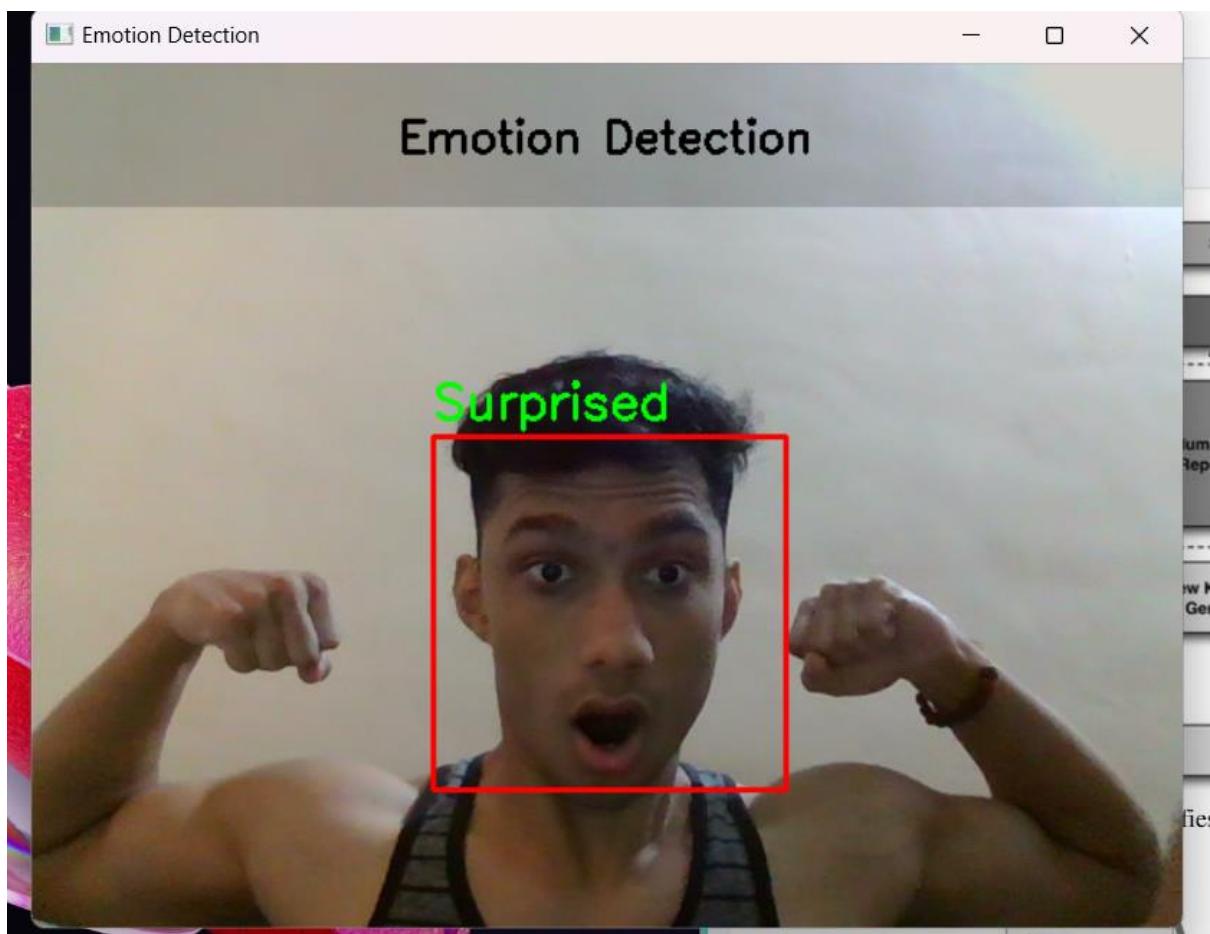
The screenshot shows a Jupyter Notebook interface. On the left, there is a file browser window titled 'dataset - JupyterLab' showing a single file 'fer2013.zip'. On the right, a code editor window titled 'Face_Emotion_Recognition_vX' contains the same Python script as the previous screenshot. A message box is displayed at the bottom of the code editor window with the following text:

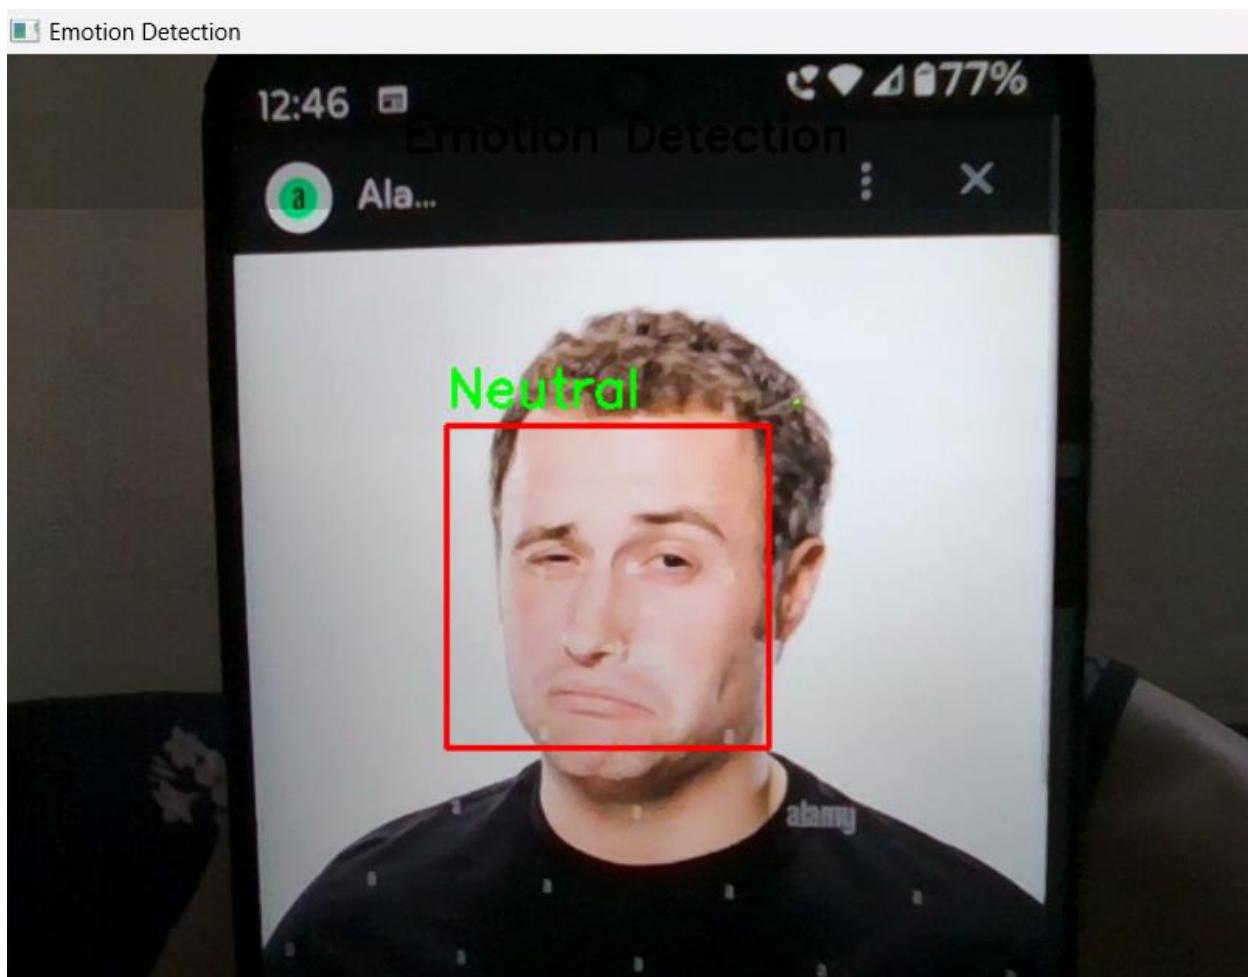
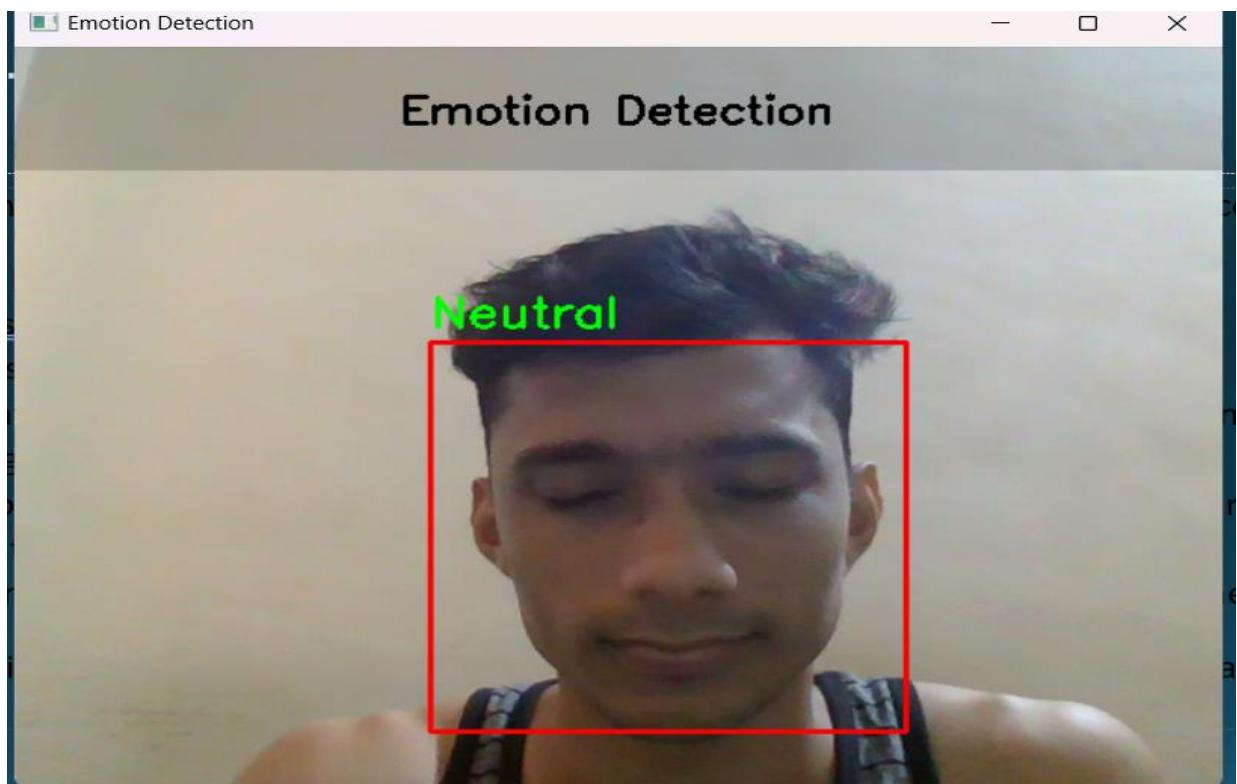
The model pretty accurately detected all the emotions beside disgust. The training data for the emotion class - disgust were pretty low as compared to others. Here are the results:

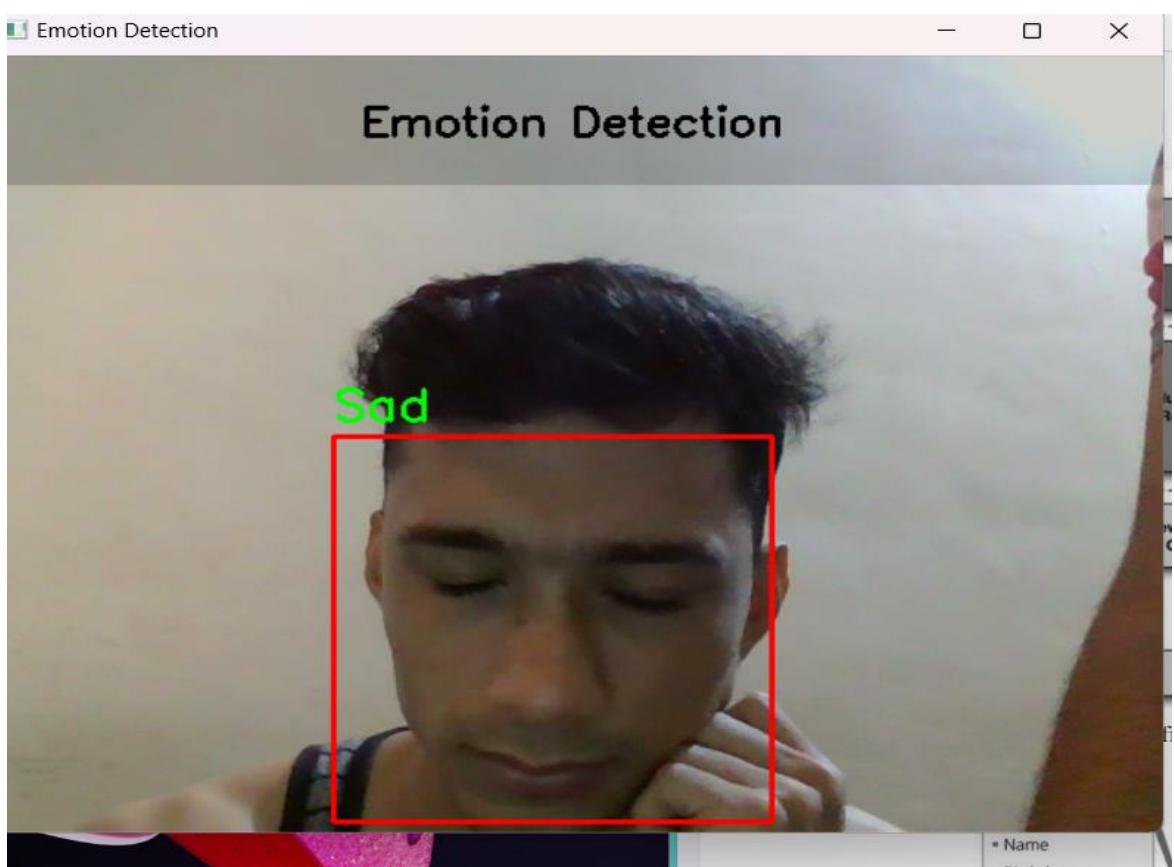
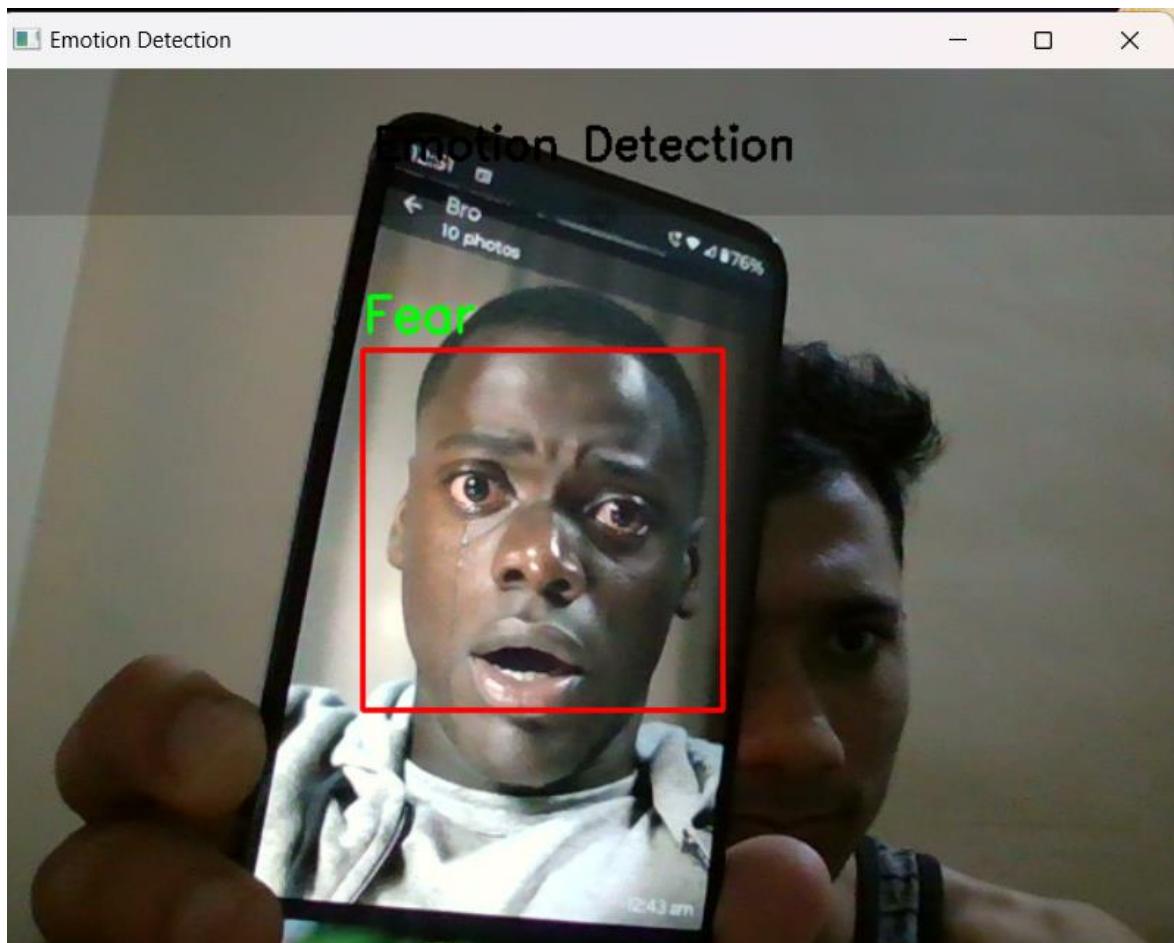
LIVE FACIAL EXPRESSION AND EMOTION RESULTS:





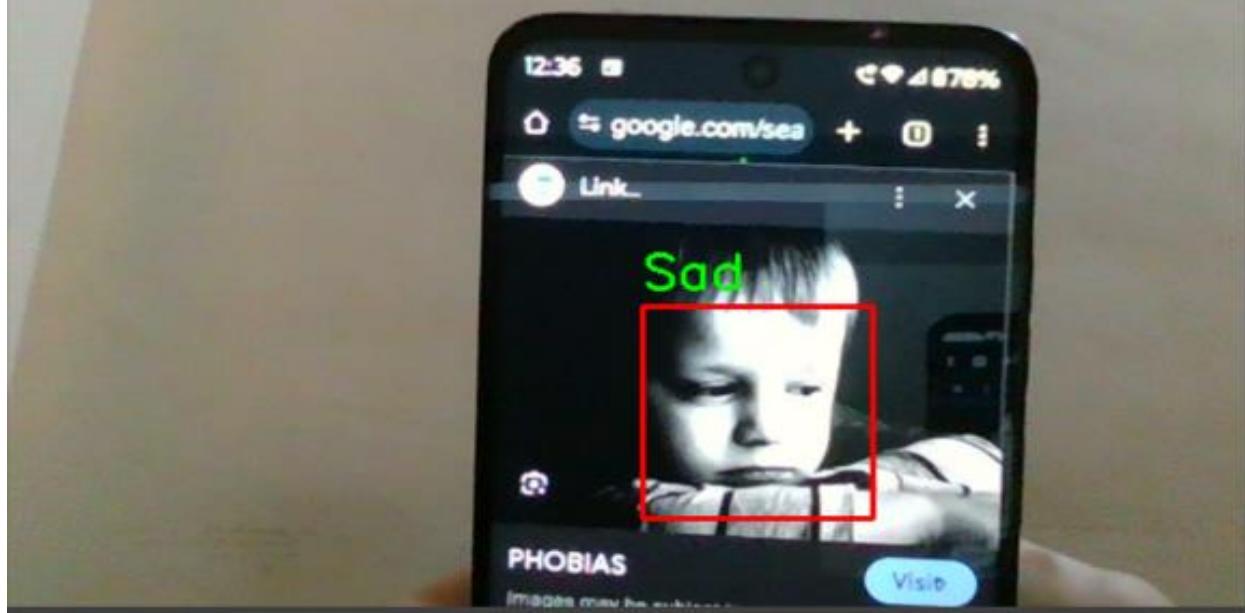






Emotion Detection

Emotion Detection



Conclusion

Some important conclusions in the project includes:

Best training result

At Epoch 30

Model training result were - accuracy: 0.5625 - loss: 1.2514 -
val_accuracy: 0.8000 - val_loss: 0.4976

-here top val_accuracy of 0.8000 was achieved

At Epoch 39

Model training result were - accuracy: 0.6406 - loss: 1.0494 -
val_accuracy: 0.2000 - val_loss: 1.1991 at epoch 39

-here top accuracy of 0.6406 was achieved

- Model Test Results:

113/113		66s 584ms/step		
		precision	recall	f1-score
	0	0.54	0.37	0.44
	1	0.45	0.09	0.15
	2	0.42	0.17	0.24
	3	0.73	0.87	0.79
	4	0.41	0.49	0.45
	5	0.68	0.72	0.70
	6	0.48	0.65	0.55
	accuracy			0.57
	macro avg	0.53	0.48	0.48
	weighted avg	0.55	0.57	0.54

The predictions on the test set are pretty great with accuracy of 0.68 which implies generalization of the model is good as compared to the training results.

(0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral).

Overall Performance:

- **Accuracy:** 0.68 (68%), indicating the model correctly predicts the emotions 68% of the time.

Summary:

The model performs best on the “Happy” class with high precision, recall, and F1-score, indicating it can accurately and consistently identify happy faces. However, it struggles with the “Disgust” classes, showing lower precision and recall. The overall accuracy of 68% suggests the model generalizes well from the training data to the test data.

Note: disgust images in the above dataset were comparatively lesser compare to happy , sad etc....